

# ETC5242Assignment

Sahinya Akila, Xinyi Cui, Pranali Angne, Janice Hsin Hsu

9/4/2021


```
library(tidyverse)
library(survival)
library(survminer)
library(kableExtra)
library(knitr)
library(ggplot2)
## Remove the line break in the file name!
churn_dat <- read_csv("https://raw.githubusercontent.com/square/pysurvival/master/pysurvival/datasets/churn.csv")
churn_dat <- churn_dat %>% filter(months_active > 0) %>% select(c(company_size, months_active, churned))

km_model <- function(time, event){
  dataset <- data_frame(time, event)

  km_data <- dataset %>%
    group_by(time, event) %>%
    summarise(died = n()) %>%
    ungroup() %>%
    mutate(risk = nrow(dataset) - accumulate(died, `+`) + died) %>%
    filter(event == 1) %>%
    mutate(probability = 1 - died/risk,
           survival = accumulate(probability, `*`))
  return(km_data %>% select(time, survival))
}


km_survive <- km_model(churn_dat$months_active, churn_dat$churned)

km_survive %>%
  ggplot(aes(time, survival)) +
  geom_step()
```



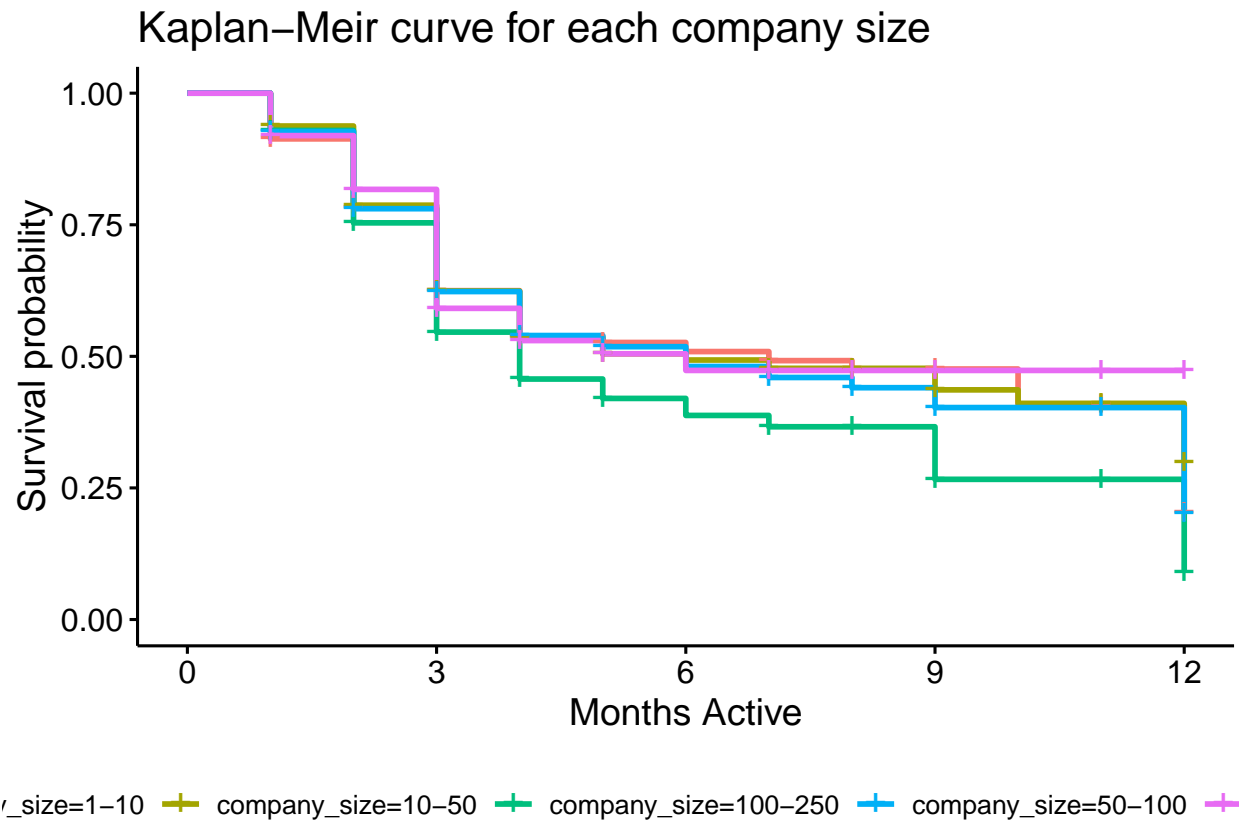
```
company_km_model <- data.frame(time = double(), survival = double(), company_size = character())
for(size in unique(churn_dat$company_size)){
  filtered <- churn_dat %>% filter(company_size == size)
  final_model <- km_model(filtered$months_active, filtered$churned) %>% mutate(company_size = size)
  company_km_model <- rbind(company_km_model, final_model)
}

company_km_model %>%
  ggplot(aes(time, survival)) +
  geom_step() +
  facet_wrap(~company_size)
```



## Question 2

Compute the Kaplan-Meier curve and use this to estimate the median churn time.



```
for (size in unique(churn_dat$company_size)){
  temp_data <- churn_dat %>% filter(company_size == size)
  name <- size
  assign(name, survfit(Surv(months_active, churned) ~ company_size, data = temp_data))
}
```

```
company_median <- data_frame(company_size = unique(churn_dat$company_size), median = c(NA, NA, NA, NA, NA))
```

```
for (i in 1:length(company_median$company_size)){
  company_median$median[i] <- median_function(get(company_median$company_size[i]))
}
```

```
company_median %>%
  knitr::kable(
    caption = "Medians for different company sizes") %>%
  kable_styling(c("hover", "striped"))
```

The table above demonstrates the median churn time estimated for different company size. - Company size of 1-10 have the highest estimated median of 7 months. - Company size of 100-250 have the lowest estimated median of 4 months. - The rest of the company sizes have the same estimated median of 5 months.

Table 1: Medians for different company sizes

company_size	median
10-50	5
100-250	4
50-100	5
1-10	7
self-employed	5

Use a non-parametric bootstrap to construct 90% confidence intervals for the median of each company size

```
bootstrapmedian <- function(df_median, df){
  bootstrap <- tibble(experiment = rep(1:1000, each = nrow(df)),
                      ind = sample(1:nrow(df), size = nrow(df)*1000, replace = TRUE),
                      timestar = df$months_active[ind],
                      churnstar = df$churned[ind])

  bias <- bootstrap %>%
    group_by(experiment) %>%
    summarise(delta = median_function(df_median) - median_function(survfit(Surv(timestar, churnstar) ~

  ci <- median_function(df_median) + quantile(bias$delta, c(0.05, 0.95))

  return(ci)
}

company_median_ci <- data_frame(company_size = unique(churn_dat$company_size), median = c(NA, NA, NA, NA, NA))

for (i in 1:length(company_median_ci$company_size)){
  ci <- bootstrapmedian(get(company_median_ci$company_size[i]), churn_dat %>% filter(company_size == company_median_ci$company_size[i]))
  company_median_ci$median[i] <- median_function(get(company_median_ci$company_size[i]))
  company_median_ci$lci[i] <- ci[1]
  company_median_ci$uci[i] = ci[2]
}

company_median_ci %>%
  knitr::kable(
    caption = "estimated mean under 90% CI" %>%
    kable_styling(c("hover", "striped"))
  )
```

\begin{table}

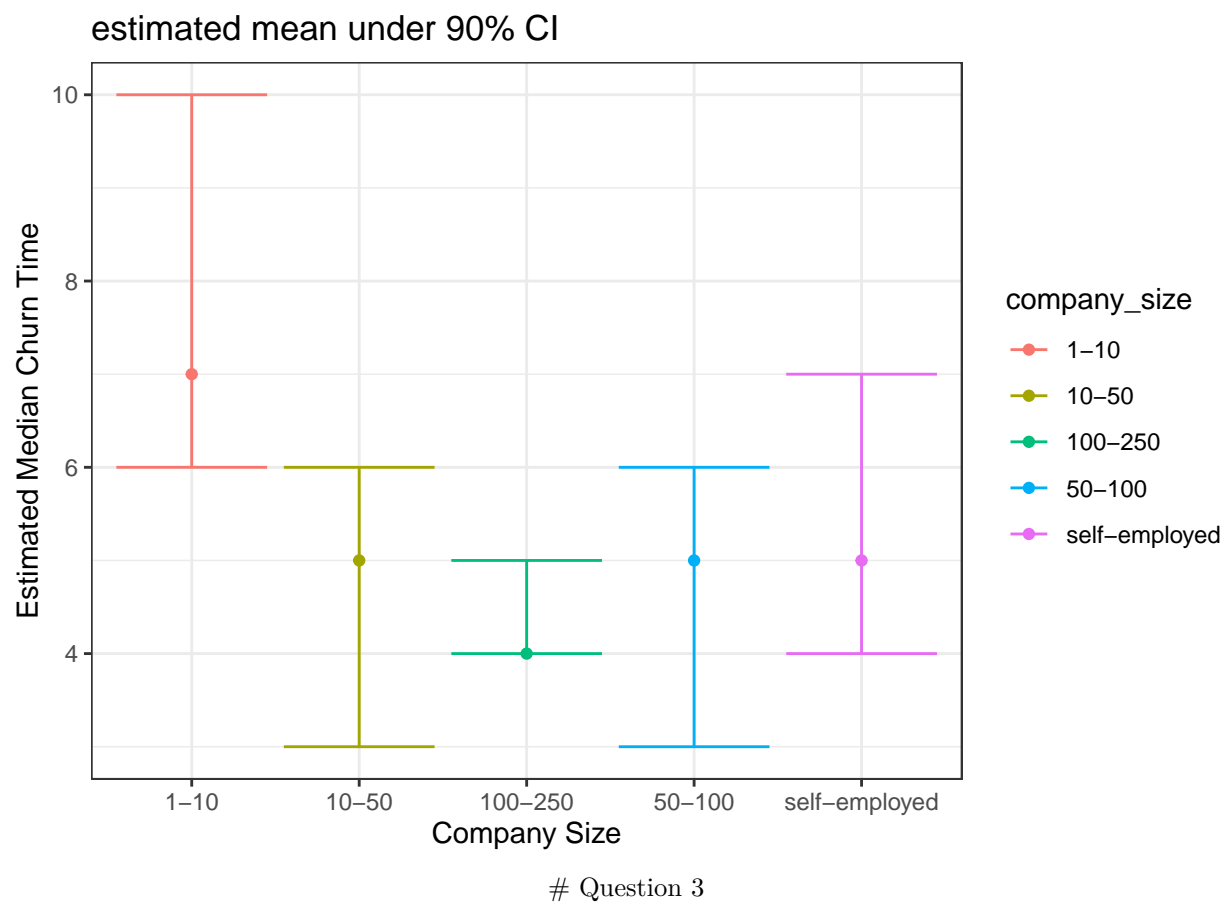
\caption{estimated mean under 90% CI}

company_size	median	lci	uci
10-50	5	3	6
100-250	4	4	5
50-100	5	3	6
1-10	7	6	10
self-employed	5	4	7

\end{table}

Make a plot that shows that estimate of the median and the corresponding confidence interval on the same axes

```
ggplot(company_median_ci,
       aes(x = company_size,
           y = median,
           colour = company_size)) +
  geom_errorbar(aes(ymax = uci, ymin = lci)) +
  geom_point() +
  theme_bw() +
  labs(x = "Company Size",
       y = "Estimated Median Churn Time",
       title = "estimated mean under 90% CI")
```



Choose company size of 50-100

```
q3_company <- churn_dat %>%
  filter(company_size == "50-100")

q3_fit <- survfit(Surv(months_active, churned) ~1,
                  data = q3_company)
```

Use a nonparametric bootstrap to re-sample the data and construct 90% confidence intervals for the survival curve at each time.

```
bootstrap_time <- tibble(experiment = rep(1:1000, each = 672),
                        ind = sample(1:672, size = 672*1000, replace = TRUE),
                        months_active = q3_company$months_active[ind],
                        churned = q3_company$churned[ind])

bias_time <- bootstrap_time %>%
  group_by(experiment) %>%
  summarise(delta = q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv)
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length
```

```
## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length

## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length

## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length

## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length

## Warning in q3_fit$surv - survfit(Surv(months_active, churned) ~ 1)$surv: longer
## object length is not a multiple of shorter object length

## 'summarise()' has grouped output by 'experiment'. You can override using the '.groups' argument.
```

```
lower <- q3_fit$surv + quantile(bias_time$delta, 0.05)
upper <- q3_fit$surv + quantile(bias_time$delta, 0.95)
```

```
Month <- c(1:11)
```

```
time_50_100_CIs <- data.frame(Month, q3_fit$surv, lower, upper) %>%
  rename("Probability" = q3_fit.surv, "Lower Confidence Interval" = lower, "Upper Confidence Interval" = upper)

kable(time_50_100_CIs,
  caption = "90% confidence intervals for the survival curve at each time for company size 50-100")
kable_styling(c("hover", "striped"))
```

```
\begin{table}

\caption{90% confidence intervals for the survival curve at each time for company size 50-100}
```

Month	Probability	Lower Confidence Interval	Upper Confidence Interval
1	0.9270833	0.8783604	0.9754078
2	0.7805394	0.7318165	0.8288638
3	0.6231333	0.5744104	0.6714577
4	0.5395180	0.4907951	0.5878424
5	0.5183604	0.4696375	0.5666848
6	0.4807375	0.4320146	0.5290619
7	0.4598358	0.4111130	0.5081603
8	0.4404062	0.3916833	0.4887306
9	0.4026571	0.3539342	0.4509815
10	0.4026571	0.3539342	0.4509815
11	0.2013285	0.1526056	0.2496530

```
\end{table}
```



Compute simultaneous coverage for the entire survival function.

Q4

log-rank test

```
q4_comp <- churn_dat %>%
  mutate(comp_hyp = case_when(company_size == "50-100" ~ 1, company_size == "100-250" ~ 2, TRUE ~ 0))

q4_comp <- q4_comp %>%
  filter(comp_hyp == 1 | comp_hyp == 2)
```

```
survdif(Surv(months_active, churned) ~ comp_hyp, data=q4_comp)
```

```
## Call:
## survdif(formula = Surv(months_active, churned) ~ comp_hyp, data = q4_comp)
##
##           N Observed Expected (O-E)^2/E (O-E)^2/V
## comp_hyp=1 672      313      332      1.14      5.26
## comp_hyp=2 240      135      116      3.27      5.26
##
## Chisq= 5.3  on 1 degrees of freedom, p= 0.02
```

```
treatment <- q4_comp$churned
outcome <- q4_comp$months_active


#Difference in means
original <- diff(tapply(outcome, treatment, mean))
mean(outcome[treatment==1]) - mean(outcome[treatment==0])
```

```
## [1] -1.896937
```

```
#Permutation test
permutation.test <- function(treatment, outcome, n){
  distribution=c()
  result=0
  for(i in 1:n){
    distribution[i]=diff(by(outcome, sample(treatment, length(treatment), FALSE), mean))
  }
  result=sum(abs(distribution) >= abs(original))/(n)
  return(list(result, distribution))
}

test1 <- permutation.test(treatment, outcome, 10000)
hist(test1[[2]], breaks=50, col='grey', main="Permutation Distribution", las=1, xlab='')
abline(v=original, lwd=3, col="red")
```

## Permutation Distribution



```
test1[[1]]
```

```
## [1] 0
```

```
#Compare to t-test
```

```
t.test(outcome~treatment)
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: outcome by treatment
```

```
## t = 13.702, df = 842.56, p-value < 2.2e-16
```

```
## alternative hypothesis: true difference in means is not equal to 0
```

```
## 95 percent confidence interval:
```

```
## 1.625195 2.168678
```

```
## sample estimates:
```

```
## mean in group 0 mean in group 1
```

```
## 4.823276 2.926339
```

## Question 5

fit a Weibull distribution to the survival data to estimate the mean and the median of the churn time for each company size

```
fit_q5 <- function(dat){  
  fit <- survreg(Surv(months_active, churned) ~ 1, data = dat, dist = "weibull" )  
  rweibull_shape <- 1 / fit$scale ## Approximately 3  
  rweibull_scale <- exp(coef(fit)) ## approximately 7  
  
  median <- rweibull_scale*log(2)^(1/rweibull_shape)  
  mean <- rweibull_scale*gamma(1+(1/rweibull_shape))  
  
  return(c(median, mean))  
}  
  
weibull <- data_frame(company_size = character(), median = double(), mean = double())  
for (size in unique(churn_dat$company_size)){  
  temp_data <- churn_dat %>% filter(company_size == size)  
  return_values <- fit_q5(temp_data)  
  weibull <- rbind(weibull, c(size, round(return_values[1], 2), round(return_values[2], 2)))  
}  
  
names(weibull) <- c("company_size", "median", "mean")  
  
kable(weibull)
```

company_size	median	mean
10-50	5.69	6.79
100-250	4.7	5.45
50-100	5.56	6.61
1-10	5.74	7
self-employed	6.23	7.92